

BlackSwanX

M&A; Intelligence Platform

Technical Overview & Capability Briefing

May 2026 • **CONFIDENTIAL**

Executive Summary

BlackSwanX is a zero-cost, fully local M&A; document intelligence platform that deploys adversarial AI agent swarms to analyse deal documents with the depth of a 10-person specialist team — in minutes, not weeks. No cloud. No API fees. No data leaving the building.

BlackSwanX runs on a standard laptop using a locally hosted 3-billion-parameter language model (llama3.2:3b via Ollama), with a pre-screening architecture that reduces LLM calls by **96%** without sacrificing coverage.

Metric	Verified Value
Annual cost	\$0 (fully local, zero API fees)
Ingestion speed	142 ms for an 850-page document
LLM call reduction	96% via pre-screening (80 vs 2,000+ calls)
Model	llama3.2:3b — 3B parameters, runs on any MacBook
Test coverage	117 automated tests, 0 warnings
Data sovereignty	Zero data leaves the machine

1. Document Processing Pipeline

1.1 Multi-Format Ingestion

Supports PDF (including scanned with embedded table extraction), DOCX, XLSX, XLS, CSV, and plain text. Financial tables embedded in PDFs — a common failure point in competing tools — are extracted and converted to pipe-delimited text for agent consumption.

1.2 Synaptic Pruning — Intelligent Pre-Screening

Before any AI agent touches a document, every text chunk is scored by M&A; signal density using **15 pattern classifiers**. A 500-page SPA is 80% boilerplate. Agents see what matters first.

Section Type	Priority Score
Representations & Warranties	95 / 100
Material Adverse Effect (MAE)	93 / 100
Indemnification	90 / 100
Net Working Capital (NWC)	90 / 100
Closing Conditions	88 / 100
Termination / Break Fee	87 / 100

Section Type	Priority Score
Purchase Price	85 / 100
Earn-Out / Contingent Consideration	85 / 100
Change of Control	85 / 100
Non-Compete / Restrictive Covenants	80 / 100
IP Ownership & Assignment	80 / 100
Financial Statements & Metrics	75 / 100
Definitions	70 / 100
Dispute Resolution / Governing Law	60 / 100
Boilerplate (Miscellaneous, Notices)	15 / 100

Performance benchmark: 142 ms to ingest and classify an 850-page document. ~80 LLM calls per standard SPA — a 96% reduction vs naive implementation.

2. Knowledge Graph Engine

2.1 Entity Extraction

The system automatically identifies and classifies six entity types:

- **Company** — acquirer, target, subsidiaries, advisors, counterparties
- **Person** — executives, signatories, key personnel, guarantors
- **Contract** — SPAs, NDAs, loan agreements, service agreements, licenses
- **Risk** — litigation, regulatory exposure, contingent liabilities, IP disputes
- **Financial Metric** — revenue, EBITDA, NWC, valuations, earn-out targets
- **EvidenceGap** — claims made in documents without supporting evidence

2.2 15 Typed Relationship Edges

Rather than generic co-occurrence links, the system establishes **15 semantically typed relationships**. The two newest — **UPDATES** and **CONTRADICTS** — are the foundation of the Liar's Drift relational versioning system (see Section 6.1):

Edge Type	Semantic Meaning
triggers_risk	Contract/clause can trigger this risk event
party_to	Entity is a named party to this contract
quantifies_risk	Financial metric quantifies a specific risk
has_risk	Entity directly carries this risk
has_revenue	Entity has this revenue/financial metric
change_of_control	Change-of-control clause connects these entities
mae_trigger	MAE clause links these entities
governed_by	Jurisdiction governing this contract/clause

Edge Type	Semantic Meaning
signatory_of	Person is a signatory on this contract
defines_value	Definition clause links term to its value
gap_about	EvidenceGap points to what is missing here
personal_service	Person-specific service obligation
co_mentioned	Fallback: entities co-appear in same clause
UPDATES ★ NEW	Newer doc supersedes this claim/metric from an earlier doc
CONTRADICTS ★ NEW	Newer doc directly contradicts this claim (sentiment flip)

2.3 Ghost Nodes — Post-Acquisition Structure Simulation

After reading deal documents, the system spawns simulated future entities representing the post-close world. These appear as dashed purple nodes in the graph:

- **[GHOST] NewCo** — the merged/combined entity
- **[GHOST] Escrow** — escrow holdback structure
- **[GHOST] NWC Peg** — working capital target entity
- **[GHOST] Earn-Out** — deferred consideration structure
- **[GHOST] Integration Gap** — identified integration risk areas

This gives deal teams a visual map of the deal structure before closing — not just what exists today, but what the documents imply will exist tomorrow. No competing platform does this.

2.4 Graph Visualisation

Rendered as an interactive Connected Papers-style D3 force graph. Node size is proportional to mention frequency. Edge colour indicates severity (red = critical, yellow = high). Click any node to focus its network. **Live KG state: 90 nodes · 534 typed edges · 12 hot nodes · avg degree 11.9.**

3. Agent Swarms

3.1 Adversarial Pollination — Predator vs Prey

The platform's most differentiated capability. Two opposing swarms run simultaneously over every clause in the document.

Predator Swarm (Buyer's Counsel Perspective)

- MAC/MAE triggers — conditions that could allow the seller to walk
- Indemnification leaks — uncapped or unlimited seller obligations
- Change-of-control provisions requiring third-party consent
- Unconditional representations impossible to satisfy
- Earn-out risks and post-closing payment obligations
- IP ownership gaps or license termination triggers
- Key-man dependencies that collapse value post-close
- Assignment restrictions blocking integration

Prey Swarm (Seller's Counsel Perspective)

- Liability caps limiting seller's total exposure

- Knowledge qualifiers ('to seller's knowledge') limiting liability scope
- Basket/deductible thresholds — minimum before indemnity kicks in
- Short survival periods — limits how long seller can be sued post-close
- MAC carve-outs protecting seller from MAE claims
- Materiality qualifiers on representations and specific indemnity exclusions

The Dissonance Map

When both swarms flag the same clause, a DISSONANCE signal fires with super-linear amplification. These are the exact paragraphs where buyer's and seller's counsel will fight in negotiation:

- **Red — DISSONANCE:** both swarms clashed here (highest deal-break risk)
- **Blue — PREDATOR_SCENT:** buyer risk flagged only
- **Green — PREY_SCENT:** seller protection flagged only

No competing platform produces an adversarial clause-level conflict map.

3.2 Leukocyte Agent — Silent Deal-Killer Scanner

Named after white blood cells. Scans for clauses that surface late in diligence and collapse timelines:

- Change-of-control provisions in third-party contracts (vendor, customer, landlord consent requirements)
- Debt acceleration triggers — does the acquisition trigger immediate loan repayment?
- Cross-default clauses — does one default cascade to all debt instruments?
- IP encumbrances — is key IP pledged as collateral to a lender?

3.3 NWC Arbitrator — Working Capital Dispute Prevention

NWC adjustments are the #1 source of post-close litigation in M&A.; This agent reads the NWC peg definition in the SPA, reads the most recent balance sheet, calculates current NWC vs the contractual peg, and quantifies the price adjustment exposure before the LOI is signed.

3.4 PMI Harmonizer — Post-Merger Integration Accounting

Maps the target company's chart of accounts to the acquirer's automatically. When two companies merge, their accounting systems differ (DATEV vs QuickBooks vs SAP). This eliminates months of manual finance team work.

3.5 Consensus Jury — Entity Resolution

When the same entity appears under multiple names across documents, three sub-agents debate:

- **Assassin** — argues they are different entities (finds distinguishing evidence)
- **Defender** — argues they are the same (finds corroborating evidence)
- **Judge** — delivers verdict: merge / keep separate / flag for human audit

3.6 Time Machine — Temporal Narrative Analysis

- **Liar's Drift detection** — flags when the narrative changes between the CIM (marketing document) and disclosure schedules. Quotes both versions side by side. PE buyers identify this as their single most-wanted feature.
- **Ghost Asset detection** — identifies assets mentioned in older documents that disappear from newer ones without explanation.
- **Value Bridge accuracy** — tracks how often the company's own historical projections matched actual results. A direct measure of management credibility.

3.7 War Room Synthesis

Aggregates all agent outputs into one executive report: Red Flag Score (0–100), traffic light verdict (green / amber / red), top risks ranked by severity with source citations and page numbers, and recommended next steps.

4. Signal Pheromone System

Inspired by biological stigmergy (ant colony signalling), every agent finding deposits a typed pheromone on the relevant entity. This creates a living heat map where risk surfaces automatically — errors find the user, not the other way around. Only nodes with intensity ≥ 0.8 across signal-bearing types (Risk, EvidenceGap, FinancialMetric, Company, Person) display a visual ring on the dashboard — routine nodes render minimally.

Pheromone Type	Triggered When
RISK_PULSE	Risk node detected — pulls related financial nodes
ANOMALY_SCENT	Inconsistency flagged — pulls related claims
DISSONANCE	Predator + Prey both flagged this clause
CHAIN_TRIGGER	Implied cascade dependency detected
DRIFT_MARKER	Narrative drift on this entity
ABSENCE_FLAG	Expected topic is missing from documents
OPPORTUNITY_BLOOM	Positive finding (favourable term, seller protection)
AUDIT_BEACON	Flagged for mandatory human review

Super-linear amplification: When 3+ agents flag the same entity independently, intensities compound — combined = sum $\times$ 1.3. High-intensity entities glow on the dashboard; cold entities render minimally.

5. Semantic Intelligence

5.1 Dual Search Mode

- **BM25 keyword search** — fast, no model inference required
- **Semantic embedding search** — finds meaning, not just keywords. 'MAC trigger' finds 'material adverse event' clauses even without exact phrase match.

5.2 Liar's Drift (Semantic Drift Detection)

Compares language and claims across documents from the same deal. Flags when the CIM describes the business differently from the SPA or disclosure schedules — and quotes both versions side by side. Private equity buyers consistently identify this as their single most-wanted due diligence feature.

5.3 Absence Detection

Identifies what is missing. If a deal of this size should have audited financials but doesn't, that absence is flagged. If there's no IP assignment agreement but the company has a software product, that gap is surfaced as an EvidenceGap node in the knowledge graph.

5.4 Chain Reaction Engine

Maps dependency cascades. If Contract A has a CoC clause, and Contract A cross-defaults to Contract B which cross-defaults to Contract C, the system traces the full chain and visualises which single triggering event could

cascade through the entire deal structure. The DFS Cascade Tracer (`trace_cascade`) performs depth-first traversal from any node — following consequence edges (`triggers_risk`, `mae_trigger`, `governed_by`, `change_of_control`, `gap_about`, `party_to`) up to 6 hops deep. Each node in the chain receives a CHAIN_TRIGGER pheromone deposit with intensity decaying by depth, making cascade hotspots visible on the graph immediately. API: GET `/api/ma/kg/cascade/{node_id}?max_depth=6`

6. Supermemory Architecture Upgrades

Four architectural patterns drawn from state-of-the-art memory system design have been integrated into BlackSwanX — all running locally with zero additional dependencies.

6.1 Liar's Drift — Relational Versioning (Graph Mutability)

Previously, Liar's Drift was detected by comparing document text. Now the Knowledge Graph itself is a **living timeline**. When an SPA is ingested after a CIM:

- Every FinancialMetric and Claim node is compared by **canonical name** across documents
- If the value or narrative changed, a red directional **UPDATES** edge is drawn: *new node* → *old node*
- If sentiment flipped (positive → negative), a **CONTRADICTS** edge is drawn instead
- The old node's **is_latest** flag flips to **false** and it renders as greyed-out in the graph
- **drift_generation** counter increments on the new node — generation 0 is the original claim, generation 2 means it has been revised twice across the deal cycle

API: POST `/api/ma/kg/stamp-drift/{doc_id}` — call after every upload. GET `/api/ma/kg/drift-timeline/{entity_name}` — returns the full mutation history as a structured JSON timeline for frontend rendering. The graph is now a living record of how the seller's narrative mutated from teaser to SPA.

6.2 Deal Team Persona — Investment Thesis Cache

Inspired by Supermemory's 50 ms user-profile cache. The PE firm or lead lawyer enters their firm's investment mandate **once**. It is stored in a local SQLite table and loaded into RAM at startup — retrieval is effectively instantaneous.

Mandate Field	Example Value
<code>firm_name</code>	Redwood Capital Partners
<code>no_go_clauses</code>	Uncapped indemnities, unlimited IP warranties
<code>max_churn_pct</code>	5% annual customer churn maximum
<code>jurisdiction</code>	Delaware law required
<code>indemnity_cap</code>	Must be capped at 30% of purchase price
<code>max_leverage</code>	4x EBITDA maximum
<code>sector_focus</code>	B2B SaaS, enterprise software
<code>custom_criteria</code>	No single customer > 20% of revenue

This mandate is injected directly into the **Predator Swarm's system prompt**. Agents no longer hunt for generic risks — they hunt specifically for violations of this firm's mandate. Every violation is prefixed [**MANDATE VIOLATION**] and assigned severity=CRITICAL.

API: GET/POST `/api/ma/deal-profile` — set once, stored permanently in the local DB, injected into every subsequent agent call automatically.

6.3 Hybrid Recall — Subgraph Triple Injection

Supermemory's architecture searches atomic memories first, then fetches raw source text — giving the LLM pre-structured context before it reads messy prose. BlackSwanX implements the equivalent pattern for M&A documents:

- Before each agent LLM call, entity names are extracted from the chunk batch
- The live Knowledge Graph is queried for all edges touching those entities
- Top edges are formatted as compact **(Subject → predicate → Object)** triples
- These triples are **prepended to the LLM user message** — before the raw chunk text

Example triple block prepended to llama3.2:3b: [KG Context — structured reasoning shortcuts] Acme Corp --[party_to]--> Share Purchase Agreement (weight: 0.95) Revenue €42M --[UPDATES]--> Revenue €38M [SUPERSEDED] (Liar's Drift) John Smith --[signatory_of]--> Share Purchase Agreement (weight: 0.90) The 3B model receives structured facts first, then raw text. This dramatically reduces the reasoning load on a small model — it no longer has to infer relationships from dense legal prose.

6.4 Long-Term Memory — The Never-Forget Brain

The three upgrades above operate within a single deal. Upgrade 4 operates **across all deals, permanently**. Every fact the system learns from any document is stored in a long-term memory layer and strengthened using the **SM-2 spaced repetition algorithm** — the same algorithm used in SuperMemo and Anki. The more times a fact is re-confirmed across different documents and deals, the stronger its memory trace becomes. Once a fact is strong enough, it is promoted to **permanent memory** and never forgotten.

Three Memory Tiers

Memory Tier	Description
Working (new)	Facts from the current document only — session-level
Episodic (strengthening)	Facts seen in 2+ documents for this company/deal
Semantic (PERMANENT)	Strength ≥ 0.75 across 3+ docs in 2+ deals — never deleted

SM-2 Spaced Repetition — How Strength Grows

Each time a fact is re-encountered in a new document, its SM-2 counter fires:

Repetition	Strength	Tier
1st encounter	0.33	Working
2nd document	0.50	Episodic
3rd document	0.60	Episodic
4th document	0.67	Episodic
5th document	0.71	Episodic
6th — 3 deals	0.75 → PERMANENT	Semantic (never forgotten)

What the System Remembers Across Deals

- **Company memory:** "We analysed Acme Corp 18 months ago — their churn was 12.3%, not the 5% claimed in the CIM. This is now permanent memory."
- **Advisor fingerprint:** "Skadden always inserts uncapped IP warranties in SaaS SPAs. We have seen it in 4 deals. Flag immediately on first sight."

- **Risk pattern:** "Revenue growth > 30% + customer concentration > 50% correlated with post-close NWC dispute in 3 prior deals."
- **Management track record:** "This CEO ran TechCo before Acme. Their projections were consistently 20% optimistic in the prior deal."

Context injection order for every agent call: 1. ★ PERMANENT LTM memories (cross-deal facts, highest signal) 2. ♦ KG triples (current-document structure) 3. Raw chunk text (the document itself) The 3B model reads the most-consolidated knowledge first and the raw text last — exactly how a senior M&A; lawyer reads a document after years of deal experience.

API Endpoint	Function
GET /api/ma/memory/stats	Total memories, permanent count, by tier and type
GET /api/ma/memory/permanent	All permanently consolidated memories
POST /api/ma/memory/recall	Recall memories for given entity names
POST /api/ma/memory/learn	Manually teach or reinforce a fact (quality 0–5)
POST /api/ma/memory/consolidate/{doc_id}	Re-consolidate any existing document
GET /api/ma/memory/review-due	Memories whose SM-2 interval has passed

7. Technical Architecture

Component	Technology
Backend	Python 3.12 / Flask
Database	SQLite (local, zero-config, zero network)
LLM (agents)	llama3.2:3b via Ollama (fully local, CPU/GPU)
LLM (fact extraction)	mistral via Ollama (fully local)
Embeddings	Ollama nomic-embed-text (local)
Search	BM25 + cosine similarity
Visualisation	D3.js v7 (force-directed graph)
Document parsing	pdfplumber, python-docx, openpyxl
Frontend	Vanilla JS + CSS (zero framework dependency)
Test suite	117 tests passing, 0 warnings

Model specification: llama3.2:3b — 3 billion parameters, quantized, runs on a standard MacBook with no GPU required. Temperature 0.1 (near-deterministic). Max 300–600 tokens per agent response. Pre-screening reduces calls to ~80 per 500-page SPA. **Ingestion and indexing completes in under 2 minutes; full 174-agent analysis runs in 5–15 minutes depending on document size.**

8. Verified Metrics

Metric	Verified Value
Typed edge relationship types	15 (incl. UPDATES + CONTRADICTS)

Metric	Verified Value
Specialised agent types	8
M&A; section classifiers	15
Supermemory upgrades implemented	4 (Liar's Drift, Persona Cache, Hybrid Recall, LTM)
LTM memory tiers	3 — Working → Episodic → Semantic (permanent)
SM-2 permanent threshold	Strength ≥ 0.75 , 3+ repetitions, 2+ deals
Ingestion + indexing speed	< 2 min for a 500-page document
Ingestion benchmark (850 pages)	142 ms text extraction + classification
Full analysis time (500-page SPA)	5–15 minutes (174 agents, fully local)
LLM calls per standard SPA	~80 (96% reduction vs naive 2,000+)
LLM models	llama3.2:3b + mistral — both fully local via Ollama
Data exposure	Zero — nothing leaves the machine
Test suite	117 passing, 0 warnings
Supported file formats	PDF, DOCX, XLSX, XLS, CSV
Knowledge graph (live)	90 nodes · 534 typed edges · 12 hot nodes · avg degree 11.9
Spider-web densifier	Leaf nodes 62 → 14 · avg degree 2.3 → 11.9 (3-pass algorithm)
DFS cascade tracer	max_depth=6 · CHAIN_TRIGGER pheromones on full chain · 27 nodes/trace
Pheromone ring threshold	≥ 0.8 intensity · signal types only · 18 active nodes (was 46)

BlackSwanX — Confidential | All metrics verified from production codebase | May 2026